

# JaSPEx: Speculative Parallel Execution of Java Applications

Ivo Anjo   João Cachopo

ESW

INESC-ID Lisboa/Instituto Superior Técnico

Setembro de 2009

# Multicore Processors



# Multicore Processors

- Are the future

# Multicore Processors

- Are the future
- Are only fully used when all cores have work

# Multicore Processors

- Are the future
- Are only fully used when all cores have work
- New processors speed up all applications

# Multicore Processors

- Are the future
- Are only fully used when all cores have work
- ~~New processors speed up all applications~~
  - Most existing code is not parallel

## Problem

## Problem

- New processors do not speed up legacy applications



## Problem

- New processors do not speed up legacy applications
- Problematic to rewrite or adapt these applications

## Problem

- New processors do not speed up legacy applications
- Problematic to rewrite or adapt these applications

## Solution

## Problem



- New processors do not speed up legacy applications
- Problematic to rewrite or adapt these applications

## Solution

- Automatic Parallelization

```
void method() {  
    doA();  
    doB();  
}
```

```
void method() {  
    doA();  
    doB();  
}
```



Problem

## Problem

- Can `doA()` and `doB()` be run in parallel?

## Problem

- Can `doA()` and `doB()` be run in parallel?
  - Shared state



## Problem

- Can `doA()` and `doB()` be run in parallel?

## Parallelizing Compilers

## Problem

- Can `doA()` and `doB()` be run in parallel?

## Parallelizing Compilers

- Only if we can prove that they can never interfere

## Problem

- Can `doA()` and `doB()` be run in parallel?

## Parallelizing Compilers

- Only if we can prove that they can never interfere

## Speculative Parallelization

## Problem

- Can `doA()` and `doB()` be run in parallel?

## Parallelizing Compilers

- Only if we can prove that they can never interfere

## Speculative Parallelization

- Maybe, let's try and we'll see how it goes

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Parallelization system does not have to prove that a parallelization is valid

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Parallelization system does not have to prove that a parallelization is valid
- Uses memory transactions, that are aborted and undone if the original execution semantics are violated

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)



# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Most TLS approaches need hardware support for transactions (HTM)

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Most TLS approaches need hardware support for transactions (HTM)
  - Limits duration

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Most TLS approaches need hardware support for transactions (HTM)
  - Limits duration
  - Limits size

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Most TLS approaches need hardware support for transactions (HTM)
  - Limits duration
  - Limits size
  - Do not work on my new PC!!

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Most TLS approaches need hardware support for transactions (HTM)
  - Limits duration
  - Limits size
  - Mainstream architectures offer no such support

# Speculative Parallelization

a.k.a. Thread-Level Speculation (TLS)

- Most TLS approaches need hardware support for transactions (HTM)
  - Limits duration
  - Limits size
  - Mainstream architectures offer no such support
- Our proposal: TLS using a Software Transactional Memory (STM)



Java Speculative Parallel Executor



## Java Speculative Parallel Executor

- TLS system for Java/JVM

## Java Speculative Parallel Executor

- TLS system for Java/JVM
- Works at the JVM bytecode level

## Java Speculative Parallel Executor

- TLS system for Java/JVM
- Works at the JVM bytecode level
- Uses the JVSTM

## Java Speculative Parallel Executor

- TLS system for Java/JVM
- Works at the JVM bytecode level
- Uses the JVSTM
  - Java Versioned Software Transactional Memory

## Java Speculative Parallel Executor

- TLS system for Java/JVM
- Works at the JVM bytecode level
- Uses the JVSTM
  - Java Versioned Software Transactional Memory
  - STM implemented in Java

## Java Speculative Parallel Executor

- TLS system for Java/JVM
- Works at the JVM bytecode level
- Uses the JVSTM
  - Java Versioned Software Transactional Memory
  - STM implemented in Java
- Also done in Java



## Static Modification Module



Static Modification Module

Runtime Control Module

## Static Modification Module

- ClassLoader that rewrites bytecode

## Runtime Control Module

## Static Modification Module

- ClassLoader that rewrites bytecode
- *Transactification*

## Runtime Control Module

## Static Modification Module

- ClassLoader that rewrites bytecode
- *Transactification*
- Prevention of nontransactional operations

## Runtime Control Module

## Static Modification Module

- ClassLoader that rewrites bytecode
- *Transactification*
- Prevention of nontransactional operations

## Runtime Control Module

- Controls speculation

# *Transactification*



- No support for memory transactions on the JVM

- No support for memory transactions on the JVM
- JVSTM is just a Java library



- No support for memory transactions on the JVM
- JVSTM is just a Java library
- Application needs to be modified to use the JVSTM

# *Transactification*

- No support for memory transactions on the JVM
- JVSTM is just a Java library
- Application needs to be modified to use the JVSTM  
⇒ “Transactification”



```
jvstm.VBox<E>
```

```
jvstm.VBox<E>
```

```
jvstm.Transaction
```

```
jvstm.VBox<E>
```

- Container that keeps the version history of a memory position

```
jvstm.Transaction
```

`jvstm.VBox<E>`

- Container that keeps the version history of a memory position
  - Fields (class or instance)

`jvstm.Transaction`

`jvstm.VBox<E>`

- Container that keeps the version history of a memory position
  - Fields (class or instance)
  - Array positions

`jvstm.Transaction`



`jvstm.VBox<E>`

- Container that keeps the version history of a memory position
  - Fields (class or instance)
  - Array positions
  - Local variables

`jvstm.Transaction`

`jvstm.VBox<E>`

- Container that keeps the version history of a memory position
- `E get()`

`jvstm.Transaction`

`jvstm.VBox<E>`

- Container that keeps the version history of a memory position
- `E get()`
- `void put(E newE)`

`jvstm.Transaction`

`jvstm.VBox<E>`

- Container that keeps the version history of a memory position
- `E get()`
- `void put(E newE)`

`jvstm.Transaction`

- `Transaction begin()`

## `jvstm.VBox<E>`

- Container that keeps the version history of a memory position
- `E get()`
- `void put(E newE)`

## `jvstm.Transaction`

- `Transaction begin()`
- `void commit()`

## `jvstm.VBox<E>`

- Container that keeps the version history of a memory position
- `E get()`
- `void put(E newE)`

## `jvstm.Transaction`

- `Transaction begin()`
- `void commit()`
- `void abort()`

# *Transactification*



```
public class Person {  
    public String name;  
  
    public Person(String name) {  
        this.name = name;  
    }  
  
    public String name() {  
        return name;  
    }  
}
```



# Transactification

```
public class Person {  
    public VBox<String> $box_name;  
  
    public Person(String name) {  
        this.name = name;  
    }  
  
    public String name() {  
        return name;  
    }  
}
```

- Replace the original fields with VBoxes

```
public class Person {  
    private VBox<String> $box_name;  
  
    public Person(String name) {  
        this.name = name;  
    }  
  
    public String name() {  
        return name;  
    }  
}
```

- Replace the original fields with VBoxes

# Transactification

```
public class Person {  
    private VBox<String> $box_name = new VBox<String>();  
  
    public Person(String name) {  
        this.name = name;  
    }  
  
    public String name() {  
        return name;  
    }  
}
```

- Add VBox initializations

# Transactification

```
public class Person {  
    private VBox<String> $box_name = new VBox<String>();  
  
    public Person(String name) {  
        this.name = name;  
    }  
  
    public String name() {  
        return name;  
    }  
  
    String $box_name_get() { return $box_name.get(); }  
    void $box_name_put(String name) { $box_name.put(name); }  
}
```

- Add get and put methods, that control access to the VBox

# Transactification

```
public class Person {  
    private VBox<String> $box_name = new VBox<String>();  
  
    public Person(String name) {  
        this.name = name;  
    }  
  
    public String name() {  
        return name;  
    }  
  
    public String $box_name_get() { return $box_name.get(); }  
    public void $box_name_put(String name) { $box_name.put(name); }  
}
```

- Add get and put methods, that control access to the VBox

# Transactification

```
public class Person {  
    private VBox<String> $box_name = new VBox<String>();  
  
    public Person(String name) {  
        $box_name_put(name);  
    }  
  
    public String name() {  
        return $box_name_get();  
    }  
  
    public String $box_name_get() { return $box_name.get(); }  
    public void $box_name_put(String name) { $box_name.put(name); }  
}
```

- Replace accesses to the original fields with calls to the get and put methods

# Transactification Problems



# Transactification Problems

- native methods



# Transactification Problems

- native methods
- The Sun JVM does not allow loading modified versions of classes in the `java.*` package

# Transactification Problems

- native methods
- The Sun JVM does not allow loading modified versions of classes in the `java.*` package
- ...

# Transactification Problems

- native methods
- The Sun JVM does not allow loading modified versions of classes in the `java.*` package
- ...

⇒ Nontransactional operations need to be detected and handled

# Prevention of nontransactional operations



# Prevention of nontransactional operations

Idea

# Prevention of nontransactional operations

## Idea

- Before executing each nontransactional operation, add a call to the framework

# Prevention of nontransactional operations

## Idea

- Before executing each nontransactional operation, add a call to the framework

```
...  
System.out.println("Hello World!");  
...
```

# Prevention of nontransactional operations

## Idea

- Before executing each nontransactional operation, add a call to the framework

```
...  
SpeculationControl.nonTransactionalActionAttempted();  
System.out.println("Hello World!");  
...
```



# Speculation



# Speculation

- Speculative execution of methods

# Speculation

- Speculative execution of methods
- When a method is invoked, some of the methods it invokes may be run speculatively

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,
```

```
public static int fib(int n) {
    if (n <= 1) return n;
    return fib(n-1) + fib(n-2);
}
```

```
public static int fib(int n) {
    SpeculationId specId =
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,
            new Object[] { new Object[] { n-1 },
                new Object[] { n-2 } });
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } }));  
    if (n <= 1) {
```



```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } }));  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
    }  
}
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } });  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
        return n;  
    }  
}
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } });  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
        return n;  
    }  
    Future f0 = SpeculationControl.getResult(specId, ID0);
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } }));  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
        return n;  
    }  
    Future f0 = SpeculationControl.getResult(specId, ID0);  
    Future f1 = SpeculationControl.getResult(specId, ID1);
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } });  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
        return n;  
    }  
    Future f0 = SpeculationControl.getResult(specId, ID0);  
    Future f1 = SpeculationControl.getResult(specId, ID1);  
    int temp = f0.get() + f1.get();  
}
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } });  
  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
        return n;  
    }  
  
    Future f0 = SpeculationControl.getResult(specId, ID0);  
    Future f1 = SpeculationControl.getResult(specId, ID1);  
    int temp = f0.get() + f1.get();  
    SpeculationControl.exitPointReached(specId);  
}
```

```
public static int fib(int n) {  
    if (n <= 1) return n;  
    return fib(n-1) + fib(n-2);  
}
```

```
public static int fib(int n) {  
    SpeculationId specId =  
        SpeculationControl.entryPointReached(ENTRY_POINT_ID,  
            new Object[] { new Object[] { n-1 },  
                           new Object[] { n-2 } });  
  
    if (n <= 1) {  
        SpeculationControl.exitPointReached(specId);  
        return n;  
    }  
  
    Future f0 = SpeculationControl.getResult(specId, ID0);  
    Future f1 = SpeculationControl.getResult(specId, ID1);  
    int temp = f0.get() + f1.get();  
    SpeculationControl.exitPointReached(specId);  
    return temp;  
}
```

# Doing Speculation





# Doing Speculation

- Call to `entryPointReached(...)` generates tasks

# Doing Speculation

- Call to `entryPointReached(...)` generates tasks
  - `fib(n-1)`

# Doing Speculation

- Call to `entryPointReached(...)` generates tasks
  - `fib(n-1)`
  - `fib(n-2)`

# Doing Speculation

- Call to `entryPointReached(...)` generates tasks
  - `fib(n-1)`
  - `fib(n-2)`
- Each task is executed inside an STM transaction

# Transaction Commit



# Transaction Commit

- Tasks wait permission for commit

# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations

# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination



# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination
  - Join with child speculation task

# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination
  - Join with child speculation task
- Commit needs to guarantee the original semantics

# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination
  - Join with child speculation task
- Commit needs to guarantee the original semantics

# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination
  - Join with child speculation task
- Commit needs to guarantee the original semantics
  - ⇒ Threads commit in program-order

# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination
  - Join with child speculation task
- Commit needs to guarantee the original semantics  
⇒ Threads commit in program-order
- At each moment, only one thread is running in program-order, and it can yield the program-order to other tasks


# Transaction Commit

- Tasks wait permission for commit
  - Execute nontransactional operations
  - Termination
  - Join with child speculation task
- Commit needs to guarantee the original semantics  
⇒ Threads commit in program-order
- At each moment, only one thread is running in program-order, and it can yield the program-order to other tasks
  - Which can then try to commit

# Preliminary Experimental Results



# Preliminary Experimental Results



Time for calculating `fib(50)` with 1..8 cores.  
A `fib` version modified with a threshold was used.



# Conclusions



# Conclusions

- *Transactification* of Java/JVM applications

# Conclusions

- *Transactification* of Java/JVM applications
- *Transactification* without support from the JVM

# Conclusions

- *Transactification* of Java/JVM applications
- *Transactification* without support from the JVM
  - Hard

# Conclusions

- *Transactification* of Java/JVM applications
- *Transactification* without support from the JVM
  - Hard
  - Imposes limitations

# Conclusions

- *Transactification* of Java/JVM applications
- *Transactification* without support from the JVM
  - Hard
  - Imposes limitations
- It is possible to achieve speedups for some applications

# Future Work



- Realistic benchmarks



# Future Work

- Realistic benchmarks
- Reducing overheads

# Future Work

- Realistic benchmarks
- Reducing overheads
  - Optimize the JVSTM for our use-case

# Future Work

- Realistic benchmarks
- Reducing overheads
  - Optimize the JVSTM for our use-case
- Gather statistics

# Future Work

- Realistic benchmarks
- Reducing overheads
  - Optimize the JVSTM for our use-case
- Gather statistics
- Support *transactification* at the JVM-runtime level

Thank you!

Questions?

