

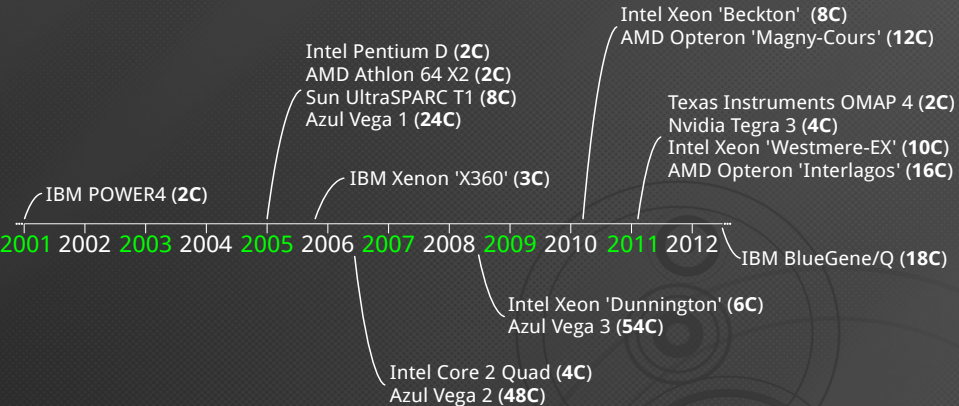
# A Software-based Method-Level Speculation Framework for the Java Platform

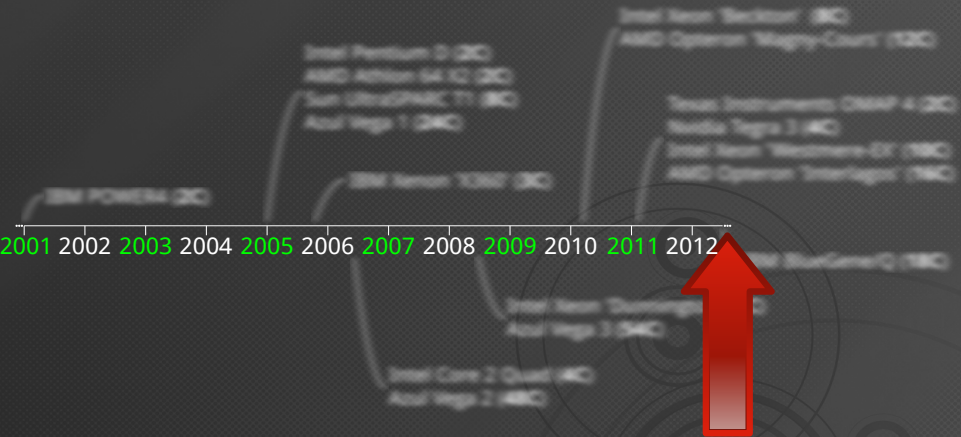
Ivo Anjo João Cachopo

ESW

INESC-ID Lisboa/Instituto Superior Técnico/Universidade Técnica de Lisboa

September 2012





Many applications still single or lightly-threaded

Not feasible to rewrite many existing applications to work in parallel



Not feasible to rewrite many existing applications to work in parallel

→ Automatic Parallelization



# Thread-Level Speculation

Employs speculation


- No need to prove that parallelization is valid



# Thread-Level Speculation

Employs speculation

- No need to prove that parallelization is valid
- Correctness dynamically ensured at runtime



# Thread-Level Speculation

Employs speculation

- No need to prove that parallelization is valid
  - Correctness dynamically ensured at runtime
- ➔ Uses some kind of in-memory transactions to buffer and validate tentative changes



# JaSPEx-MLS

Our proposal: JaSPEx-MLS

- Software-based speculative parallelization framework for Java



# JaSPEx-MLS

Our proposal: JaSPEx-MLS

- Software-based speculative parallelization framework for Java
- Java Bytecode rewriting



# JaSPEx-MLS

Our proposal: JaSPEx-MLS

- Software-based speculative parallelization framework for Java
- Java Bytecode rewriting
- Automatic transactification



# JaSPEX-MLS


Our proposal: JaSPEX-MLS

- Software-based speculative parallelization framework for Java
- Java Bytecode rewriting
- Automatic transactification
  - STM-like API




# JaSPEX-MLS

Our proposal: JaSPEX-MLS

- Software-based speculative parallelization framework for Java
  - Java Bytecode rewriting
  - Automatic transactification
    - STM-like API
  - Handling non-transactional operations
- 


# JaSPEx-MLS

Our proposal: JaSPEx-MLS

- Software-based speculative parallelization framework for Java
  - Java Bytecode rewriting
  - Automatic transactification
    - STM-like API
  - Handling non-transactional operations
  - Method-level speculation
- 
- A decorative background graphic consisting of several overlapping circles of varying sizes and shades of gray, creating a ripple effect. The circles are centered in the lower right quadrant of the slide.

# JaSPEX-MLS

Our proposal: JaSPEX-MLS

- Software-based speculative parallelization framework for Java
  - Java Bytecode rewriting
  - Automatic transactification
    - STM-like API
  - Handling non-transactional operations
  - Method-level speculation
  - Nested speculation
- 
- A decorative background graphic consisting of several overlapping circles of varying sizes and shades of gray, creating a ripple effect. The circles are centered in the lower right quadrant of the slide.

# JaSPEX-MLS

Our proposal: JaSPEX-MLS

- Software-based speculative parallelization framework for Java
- Java Bytecode rewriting
- Automatic transactification
  - STM-like API
- Handling non-transactional operations
- Method-level speculation
- Nested speculation

...on top of the OpenJDK Java VM



# JaSPEX-MLS

Our proposal: JaSPEX-MLS

- Software-based speculative parallelization framework for Java
- Java Bytecode rewriting
- Automatic transactification
  - STM-like API
- Handling non-transactional operations
- Method-level speculation
- Nested speculation

...on top of the OpenJDK Java VM + Continuation support

# JaSPEx-MLS

Target applications

- Existing single-threaded JVM bytecode



# JaSPEx-MLS

## Target applications

- Existing single-threaded JVM bytecode
- Object-oriented / method-heavy



# JaSPEx-MLS

## Target applications

- Existing single-threaded JVM bytecode
- Object-oriented / method-heavy
  - Limited support for loops



# JaSPEx-MLS

## Target applications

- Existing single-threaded JVM bytecode
- Object-oriented / method-heavy
  - Limited support for loops
- Static analysis is hard or impossible

# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors



# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
- Just-in-time compilation



# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
- Just-in-time compilation
- Adaptive optimization





# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
- Just-in-time compilation
- Adaptive optimization
- Advanced garbage collection




# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
- Just-in-time compilation
- Adaptive optimization
- Advanced garbage collection
- Java 6 support



# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
  - Just-in-time compilation
  - Adaptive optimization
  - Advanced garbage collection
  - Java 6 support
  - Optimized concurrency primitives
- 
- A decorative background graphic consisting of several overlapping circles of varying sizes and shades of gray, creating a ripple effect. The circles are centered in the lower right quadrant of the slide.

# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
  - Just-in-time compilation
  - Adaptive optimization
  - Advanced garbage collection
  - Java 6 support
  - Optimized concurrency primitives
- ➔ JaSPEx-MLS works on commonly available hardware

# Why use OpenJDK?

- Modern open-source VM by Oracle and community contributors
  - Just-in-time compilation
  - Adaptive optimization
  - Advanced garbage collection
  - Java 6 support
  - Optimized concurrency primitives
- 
- ➔ JaSPEX-MLS works on commonly available hardware
  - ➔ Same codebase used to run regular Java applications


# Method-Level Speculation

```
int example1() {  
    int x = computeValue();  
    int y = 0;  
    for (...) y += ...;  
    return x+y;  
}
```

- Run method call in parallel with code following its return

# Method-Level Speculation



```
int example1() {  
    int x = computeValue();  
    int y = 0;  
    for (...) y += ...;  
    return x+y;  
}
```



- Run method call in parallel with code following its return

# Method-Level Speculation

```
int example1() {  
    int x = computeValue();  
    int y = 0;  
    for (...) y += ...;  
    return x+y;  
}
```



- Run method call in parallel with code following its return



# JaSPEx-MLS

## ClassLoader

- Prepares code for speculative parallelization



# JaSPEx-MLS

## ClassLoader

- Prepares code for speculative parallelization

## Runtime Orchestration Library

- Coordinates all runtime events
  - Creation, validation, commit, abort, ...

# JaSPEx-MLS Classloader

Transactification

- Intercept access to slots and arrays



# JaSPEx-MLS Classloader

## Transactification

- Intercept access to slots and arrays

```
object.field = 100l;
```

# JaSPEx-MLS Classloader

## Transactification

- Intercept access to slots and arrays

```
object.field = 100l;
```

```
Transaction.storeLong(object, field, 100l);
```

# JaSPEx-MLS Classloader

## Transactification

- Intercept access to slots and arrays

```
object.field = 100l;
```

```
Transaction.storeLong(object, field, 100l);
```

- Static and type-specific
  - Designed to be inlined by the VM

# JaSPEx-MLS Classloader

Handling Non-Transactional Operations



# JaSPEx-MLS Classloader

Handling Non-Transactional Operations

➔ Operations outside the control of the transactional system





# JaSPEx-MLS Classloader

Handling Non-Transactional Operations

- ➔ Operations outside the control of the transactional system
  - native methods



# JaSPEX-MLS Classloader

Handling Non-Transactional Operations

➔ Operations outside the control of the transactional system

- native methods
- JDK code



# JaSPEX-MLS Classloader

## Handling Non-Transactional Operations

→ Operations outside the control of the transactional system

- native methods
- JDK code

```
...
```

```
System.out.println("Hello World!");
```

```
...
```

# JaSPEX-MLS Classloader

## Handling Non-Transactional Operations

→ Operations outside the control of the transactional system

- native methods
- JDK code

```
...  
SpeculationControl.nonTransactionalActionAttempted();  
System.out.println("Hello World!");  
...
```

# JaSPEx-MLS Classloader

Modifications for MLS

- Replace method calls with call to `spawnSpeculation()`



# JaSPEx-MLS Classloader

Modifications for MLS

- Replace method calls with call to `spawnSpeculation()`
- Replace method return value with `Future`




# JaSPEX-MLS Classloader

## Modifications for MLS

- Replace method calls with call to `spawnSpeculation()`
- Replace method return value with `Future`

```
int mlsExample1() {  
    return compute(0) + compute(1);  
}
```




```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

# JaSPEX-MLS Classloader

## Modifications for MLS

- Replace method calls with call to `spawnSpeculation()`
- Replace method return value with `Future`

```
int mlsExample1() {  
    return compute(0) + compute(1);  
}
```



```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

- Usually needs control flow analysis and code duplication



# JaSPEx-MLS Classloader

STM support for Futures

- Allow Futures to be “written” to memory locations
- Delay obtaining value until it is needed (read or commit)

# JaSPEx-MLS Classloader

STM support for Futures

*// Original Method*

```
void doCompute(Object[] results) {  
    for (int i = 0; i < results.length; i++) {  
        results[i] = compute(i);  
    }  
}
```

# JaSPEx-MLS Classloader

STM support for Futures

*// Original Method*

```
void doCompute(Object[] results) {  
    for (int i = 0; i < results.length; i++) {  
        results[i] = compute(i);  
    }  
}
```

*// \*Attempted\* parallelization*

```
void doCompute(Object[] results) {  
    for (int i = 0; i < results.length; i++) {  
        Future f0 = spawnSpeculation(...); // compute(i)  
        TM.storeObjectArray(results, i, f0.get());  
    }  
}
```

# JaSPEX-MLS Classloader

... let's try this again

*// Original Method*

```
void doCompute(Object[] results) {  
    for (int i = 0; i < results.length; i++) {  
        results[i] = compute(i);  
    }  
}
```

# JaSPEX-MLS Classloader

... let's try this again

*// Original Method*

```
void doCompute(Object[] results) {  
    for (int i = 0; i < results.length; i++) {  
        results[i] = compute(i);  
    }  
}
```

*// \*Successful\* parallelization*

```
void doCompute(Object[] results) {  
    for (int i = 0; i < results.length; i++) {  
        Future f0 = spawnSpeculation(...); // compute(i)  
        TM.storeFutureObjectArray(results, i, f0); // f0 handed  
    } // to the STM  
}
```

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation
- Submit child work task with continuation to pool



# JaSPEx-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation
- Submit child work task with continuation to pool
- Parent task:

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation
- Submit child work task with continuation to pool
- Parent task:
  - Clean current stack

# JaSPEx-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation
- Submit child work task with continuation to pool
- Parent task:
  - Clean current stack
  - Execute `compute(0)`

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation
- Submit child work task with continuation to pool
- Parent task:
  - Clean current stack
  - Execute `compute(0)`
  - Write result to child task

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

What happens inside `spawnSpeculation`?

- Check threadpool state
- Capture first-class continuation
- Submit child work task with continuation to pool
- Parent task:
  - Clean current stack
  - Execute `compute(0)`
  - Write result to child task
  - Return to poll

# JaSPEx-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

- Parent task:
  - Clean current task
  - Execute compute(0)
  - Write result to child task
  - Return to poll
- Child task:

# JaSPEx-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

- Parent task:
  - Clean current task
  - Execute compute(0)
  - Write result to child task
  - Return to poll
- Child task:
  - Start transaction

# JaSPEX-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

- Parent task:
  - Clean current task
  - Execute compute(0)
  - Write result to child task
  - Return to poll
- Child task:
  - Start transaction
  - Resume continuation



# JaSPEx-MLS Runtime Orchestration

```
int mlsExample1() {  
    Future temp1 = spawnSpeculation(...); // compute(0)  
    Future temp2 = spawnSpeculation(...); // compute(1)  
    return temp1.get() + temp2.get();  
}
```

- Parent task:
  - Clean current task
  - Execute compute(0)
  - Write result to child task
  - Return to poll
- Child task:
  - Start transaction
  - Resume continuation
  - Return from spawnSpeculation

# JaSPEx-MLS Runtime Orchestration

Committing a speculation:

- Task completed work



# JaSPEx-MLS Runtime Orchestration

Committing a speculation:


- Task completed work
- Need to obtain value from a Future, and oldest-running task



# JaSPEx-MLS Runtime Orchestration

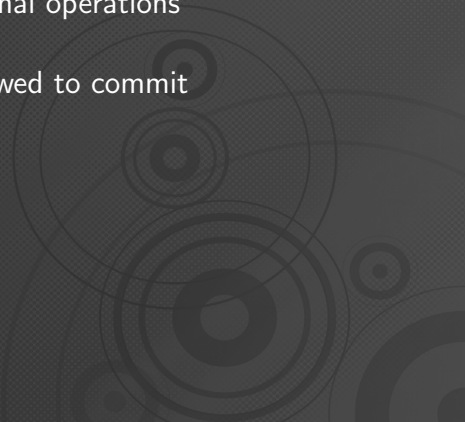
Committing a speculation:

- Task completed work
- Need to obtain value from a Future, and oldest-running task
- Tried to execute non-transactional operations



# JaSPEx-MLS Runtime Orchestration

Committing a speculation:

- Task completed work
  - Need to obtain value from a Future, and oldest-running task
  - Tried to execute non-transactional operations
  - Only oldest-running task is allowed to commit
- 
- A decorative background graphic consisting of several overlapping circles of varying sizes and shades of gray, creating a ripple effect. The circles are centered in the lower right quadrant of the slide.

# JaSPEX-MLS Runtime Orchestration

Committing a speculation:

- Task completed work
- Need to obtain value from a Future, and oldest-running task
- Tried to execute non-transactional operations
- Only oldest-running task is allowed to commit
- Wait for result from parent

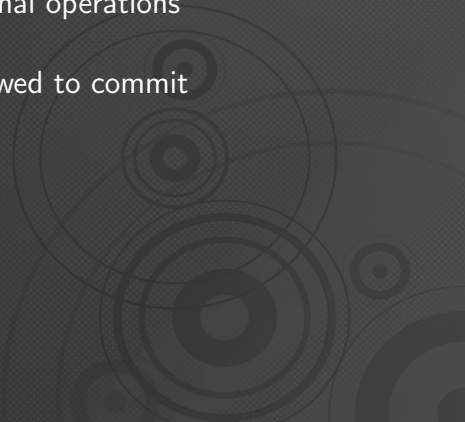
# JaSPEX-MLS Runtime Orchestration

Committing a speculation:

- Task completed work
- Need to obtain value from a Future, and oldest-running task
- Tried to execute non-transactional operations
- Only oldest-running task is allowed to commit
- Wait for result from parent
- If exception → abort

# JaSPEX-MLS Runtime Orchestration

Committing a speculation:

- Task completed work
  - Need to obtain value from a Future, and oldest-running task
  - Tried to execute non-transactional operations
  - Only oldest-running task is allowed to commit
  - Wait for result from parent
  - If exception → abort
  - Else try to commit
- 
- A decorative background graphic consisting of several overlapping circles of varying sizes and shades of gray, creating a ripple effect. The circles are centered in the lower right quadrant of the slide.



# Relaxed STM Model

Distinguish between

- Program-order mode: Reads and writes directly to memory



# Relaxed STM Model

Distinguish between

- Program-order mode: Reads and writes directly to memory
  - Not running speculatively — only one thread allowed in this mode



# Relaxed STM Model

Distinguish between

- Program-order mode: Reads and writes directly to memory
  - Not running speculatively — only one thread allowed in this mode
- Speculation mode: Read directly from memory (if not in write-set), writes go to write-set

# Relaxed STM Model

Distinguish between

- Program-order mode: Reads and writes directly to memory
  - Not running speculatively — only one thread allowed in this mode
- Speculation mode: Read directly from memory (if not in write-set), writes go to write-set
  - Inconsistent reads

# Relaxed STM Model

Distinguish between

- Program-order mode: Reads and writes directly to memory
  - Not running speculatively — only one thread allowed in this mode
- Speculation mode: Read directly from memory (if not in write-set), writes go to write-set
  - Inconsistent reads

Transaction commit is simpler, because only one transaction may commit at a time

# Relaxed STM Model

Distinguish between

- Program-order mode: Reads and writes directly to memory
  - Not running speculatively — only one thread allowed in this mode
- Speculation mode: Read directly from memory (if not in write-set), writes go to write-set
  - Inconsistent reads

Transaction commit is simpler, because only one transaction may commit at a time

➔ No synchronization needed



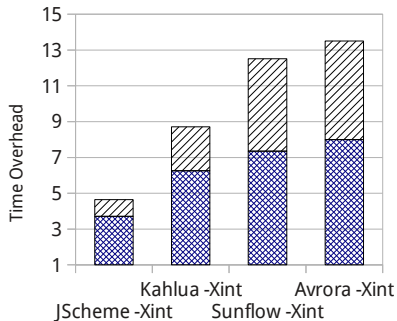
Benchmarks

# Benchmarks

- Test overheads from bytecode changes
- Early results from full system
- Test system: Intel Core i5 750, 8GB RAM, Ubuntu Linux 12.04 64-bit




# Benchmarks — Overhead with Interpreter




Transactification + spawnSpeculation

Only Transactification

# Benchmarks — Overhead with JIT




Transactification + spawnSpeculation



Only Transactification

# Benchmarks — Parallelization



- Benchmarks from Java Grande Framework
- Source was **not** modified

# Conclusions


- JaSPEX-MLS: Automatic parallelization framework for Java




# Conclusions

- JaSPEX-MLS: Automatic parallelization framework for Java
- Using custom lightweight STM and MLS


# Conclusions

- JaSPEX-MLS: Automatic parallelization framework for Java
  - Using custom lightweight STM and MLS
  - Supports nested speculation
- 
- A decorative graphic in the bottom right corner consisting of several overlapping circles of varying shades of gray, creating a ripple effect.

# Conclusions

- JaSPEx-MLS: Automatic parallelization framework for Java
  - Using custom lightweight STM and MLS
  - Supports nested speculation
  - Needs no special hardware support
- 
- A decorative graphic in the bottom right corner consisting of several overlapping circles of varying shades of gray, creating a ripple effect.

# Conclusions

- JaSPEX-MLS: Automatic parallelization framework for Java
  - Using custom lightweight STM and MLS
  - Supports nested speculation
  - Needs no special hardware support
  - On top of OpenJDK VM
- 
- A decorative graphic in the bottom right corner consisting of several overlapping circles of varying shades of gray, creating a ripple or concentric circle effect.



# Conclusions

- Integrating Futures with the STM unlocks more parallelism



# Conclusions

- Integrating Futures with the STM unlocks more parallelism
- Optimizing VM makes **\*all\*** the difference

# Conclusions

- Integrating Futures with the STM unlocks more parallelism
- Optimizing VM makes **\*all\*** the difference
  - Able to hide much of the added overhead

# Conclusions

- Integrating Futures with the STM unlocks more parallelism
- Optimizing VM makes **\*all\*** the difference
  - Able to hide much of the added overhead
- Already able to extract parallelism in some benchmarks

# Future Work

- Lift some limitations imposed by non-transactional ops




# Future Work

- Lift some limitations imposed by non-transactional ops
- Add support for profiling pass, and to save statistic information between runs



# Future Work

- Lift some limitations imposed by non-transactional ops
  - Add support for profiling pass, and to save statistic information between runs
  - Integrate task scheduler
- 

# Future Work

- Lift some limitations imposed by non-transactional ops
- Add support for profiling pass, and to save statistic information between runs
- Integrate task scheduler
- Re-use waiting threads by returning them to the thread-pool



Thank you!

The background is a dark gray with a fine halftone dot pattern. In the lower right quadrant, there are several overlapping circles of varying sizes and shades of gray, creating a ripple effect. The text "Thank you!" is centered in the upper half of the image in a white, sans-serif font.