


Speculating on top of an unmodified Java VM

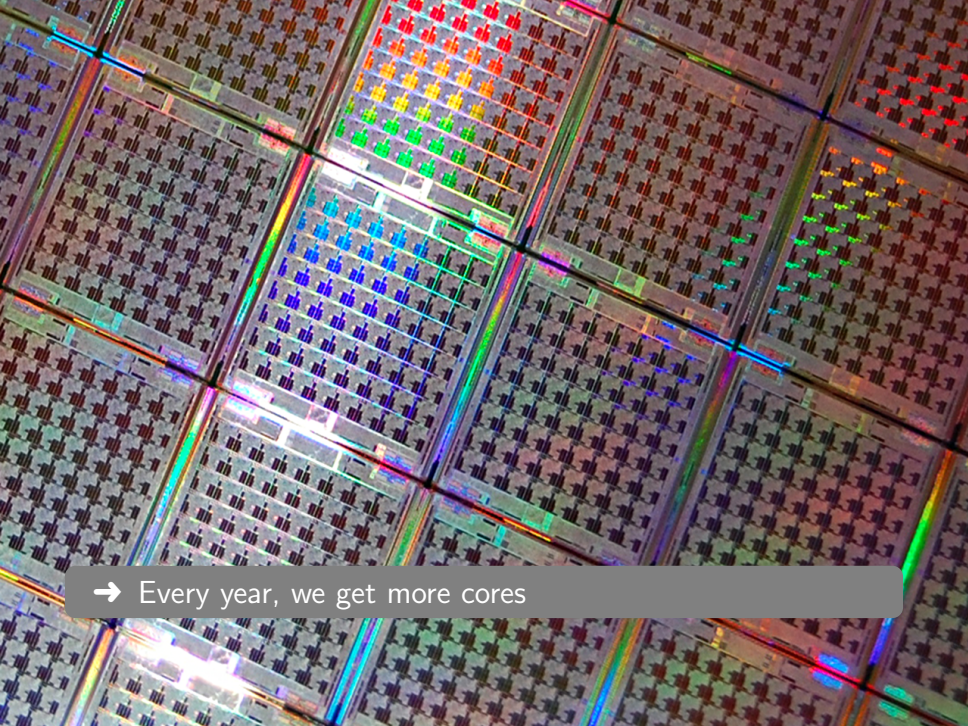
Ivo Anjo João Cachopo

ESW

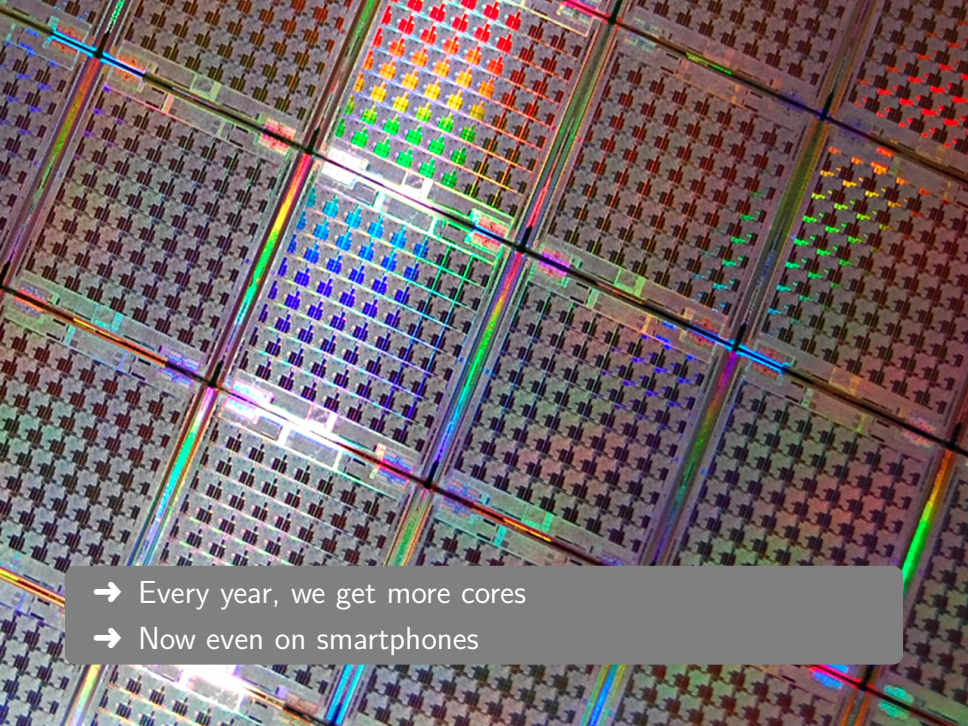
INESC-ID Lisboa/Instituto Superior Técnico

October 2011

A decorative graphic in the bottom right corner consisting of several concentric circles of varying shades of green, creating a ripple effect.



→ Every year, we get more cores



- Every year, we get more cores
- Now even on smartphones

- Many applications still not parallel

- Many applications still not parallel
- Concurrent programming still very hard

- Many applications still not parallel
- Concurrent programming still very hard
 - ➔ Automatic parallelization?

Thread-Level Speculation (TLS)

- Automatic

Thread-Level Speculation (TLS)

- Automatic
- Optimistic

Thread-Level Speculation (TLS)

- Automatic
- Optimistic
 - Uses some kind of memory transactions

```
void method() {  
    doA();  
    doB();  
    doC();  
}
```

Can doA(), doB() and doC() be executed in parallel?

TLS on the Java platform

- Modify the VM to perform TLS

TLS on the Java platform

- Modify the VM to perform TLS
- Our idea: What if we left the VM **untouched**?

TLS on the Java platform

- Modify the VM to perform TLS
- Our idea: What if we left the VM **untouched**?
 - Propose missing features as TLS-independent Java APIs

(Not) Diving into the VM

Take advantage of Java concurrency features

- Memory model

(Not) Diving into the VM

Take advantage of Java concurrency features

- Memory model
 - Programming without locks

(Not) Diving into the VM

Take advantage of Java concurrency features

- Memory model
 - Programming without locks
- Concurrency utility classes

(Not) Diving into the VM

Take advantage of Java concurrency features

- Memory model
 - Programming without locks
- Concurrency utility classes
 - Collections

(Not) Diving into the VM

Take advantage of Java concurrency features

- Memory model
 - Programming without locks
- Concurrency utility classes
 - Collections
 - Atomic operations

(Not) Diving into the VM

Take advantage of Java concurrency features

- Memory model
 - Programming without locks
- Concurrency utility classes
 - Collections
 - Atomic operations
 - Fork/Join framework

(Not) Diving into the VM

- State-of-the-art garbage collection

(Not) Diving into the VM

- State-of-the-art garbage collection
 - Oracle Java Hotspot G1 GC

(Not) Diving into the VM

- State-of-the-art garbage collection
 - Oracle Java Hotspot G1 GC
 - Azul Pauseless GC

(Not) Diving into the VM

- State-of-the-art garbage collection
 - Oracle Java Hotspot G1 GC
 - Azul Pauseless GC
- Take advantage of JIT Optimizations

(Not) Diving into the VM

- State-of-the-art garbage collection
 - Oracle Java Hotspot G1 GC
 - Azul Pauseless GC
- Take advantage of JIT Optimizations
 - Do not need to be rewritten to support TLS

(Not) Diving into the VM

- State-of-the-art garbage collection
 - Oracle Java Hotspot G1 GC
 - Azul Pauseless GC
- Take advantage of JIT Optimizations
 - Do not need to be rewritten to support TLS
- Research VM vs Production

(Not) Diving into the VM

The big disadvantage: **overheads!**

- Spawn speculation

(Not) Diving into the VM

The big disadvantage: **overheads!**

- Spawn speculation
- Transaction setup

(Not) Diving into the VM

The big disadvantage: **overheads!**

- Spawn speculation
- Transaction setup
- Transactional execution

(Not) Diving into the VM

The big disadvantage: **overheads!**

- Spawn speculation
- Transaction setup
- Transactional execution
- Speculations will need to be longer-running

(Not) Diving into the VM

The big disadvantage: **overheads!**

- Spawn speculation
- Transaction setup
- Transactional execution
- Speculations will need to be longer-running
 - Will need to support nested speculation

Challenges

Transactification

- Need to validate speculative reads
- Need to buffer speculative writes

Transactification

- Need to validate speculative reads
- Need to buffer speculative writes
- Rely on Software Transactional Memory (STM)

Transactification

- Need to validate speculative reads
- Need to buffer speculative writes
- Rely on Software Transactional Memory (STM)
 - ➔ Transactify code, intercepting field/array accesses

Transactification

Can we improve on current STMs?

- Fixed commit ordering

Transactification

Can we improve on current STMs?

- Fixed commit ordering
- Relax isolation

Transactification

Can we improve on current STMs?

- Fixed commit ordering
- Relax isolation
- Concurrent nesting

Transactification

Can we improve on current STMs?

- Fixed commit ordering
- Relax isolation
- Concurrent nesting
- Non-conservative/dynamic transactification

Non-Transactional Operations

Operations that change state outside STM control

Non-Transactional Operations

Operations that change state outside STM control

Strategies

- Static identification of native calls

Non-Transactional Operations

Operations that change state outside STM control

Strategies

- Static identification of native calls
- Runtime detection and prevention

Non-Transactional Operations

Lifting restrictions

- Extend transactional system

Non-Transactional Operations

Lifting restrictions

- Extend transactional system
 - Transactional IO

Non-Transactional Operations

Lifting restrictions

- Extend transactional system
 - Transactional IO
- Identify harmless native calls

Stack Manipulation

Saving/Reading → Continuations

Stack Manipulation

Saving/Reading → Continuations

- Bytecode modification

Stack Manipulation

Saving/Reading → Continuations

- Bytecode modification
- OpenJDK patches

Threads & Scheduling

- No control over scheduling

Threads & Scheduling

- No control over scheduling
- Thread pools

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting
 - Can we design a better pool?

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting
 - Can we design a better pool?
- Control Issues

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting
 - Can we design a better pool?
- Control Issues
 - Loops caused by inconsistent reads

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting
 - Can we design a better pool?
- Control Issues
 - Loops caused by inconsistent reads
 - Cannot kill threads

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting
 - Can we design a better pool?
- Control Issues
 - Loops caused by inconsistent reads
 - Cannot kill threads
 - User code can catch any exception

Threads & Scheduling

- No control over scheduling
- Thread pools
 - Issues with waiting
 - Can we design a better pool?
- Control Issues
 - Loops caused by inconsistent reads
 - Cannot kill threads
 - User code can catch any exception
- Green threads would be helpful

In Conclusion

- Most things needed for TLS can be done on top of the VM

In Conclusion

- Most things needed for TLS can be done on top of the VM
- Extra APIs would be helpful, but need not be TLS-specific

In Conclusion

- Most things needed for TLS can be done on top of the VM
- Extra APIs would be helpful, but need not be TLS-specific
 - ➔ Continuations, more control over threads, TM support, ...

Thank you!



